KMS vs. NYPA Enterprise Knowledge Platform

Where we stand against the EKP requirements, and how we close the gap

Goal: meet every EKP requirement of NYPA



Enterprise Knowledge Platform

NYPA wants one platform that gives employees unified search, navigation and access to knowledge across key repositories, with AI answers grounded in trusted sources.

32 requirements in RFI

1

Where we are today

Every EKP requirement rated against the actual codebase

2

What's missing

The partial and unbuilt requirements, and the biggest risks

3

How we reach 32 / 32

Workstreams, sequencing and quick wins to close every gap

How we assessed KMS

32 requirements

Every requirement in RFI grouped into 9 areas

Code-level check

API routes, database schema, ingestion pipeline and UI, not the PRD status column

Rated

Each requirement marked Met, Partial or Gap, with the evidence and what's missing

Rating scale

Met — Built and working

Partial — Foundation exists, not enough yet

Gap — Not built

Most requirements are partial: the core is built

8 — 25%

Met

Built and working

18 — 56%

Partial

Exists, needs more depth

6 — 19%

Gap

Not built yet



32 EKP requirements

In almost every partial case the working mechanism exists. What's missing is enterprise depth: configurable governance, analytics, managed taxonomy, more formats and fuller M365 integration.

Where KMS is strong and where it's weak

EKP area	Met	Partial	Gap
Unified access, search & navigation	–	4	–
Personalization	1	1	–
Conversational Q&A	3	1	–
Content coverage	–	2	2
Enrichment & semantic layer	1	–	2
Ingestion & automated metadata	2	2	–
Governance & lifecycle	1	4	1
Analytics & dashboards	–	3	1
M365 integration	–	1	–



Strongest

- Conversational Q&A
- Ingestion & classification
- Knowledge graph



Weakest

- Content coverage (formats)
- Taxonomy & semantic layer
- Analytics & dashboards

8 requirements are fully met today

#7

Natural-language answers

Org-wide chat answers questions across all documents

#8

Cited, permission-aware

Every answer cites sources; roles applied before generation

#9

Conversational follow-ups

Keeps conversation and topic context for refinement

#17

Graph-enriched discovery

Knowledge graph powers search, navigation and chat

#6

Recommendations & trending

Personal suggestions and department trending

#18

Drag-and-drop upload

Batch upload with category and department

#20

Auto-tagging with review

Low-confidence items go to a Needs Review queue

#27

Duplicate detection

Exact and near-duplicates flagged, plus contradictions

We've built more than the EKP asks for

Expert Discovery
Who knows what, inferred from real activity

Decision Intelligence
Extracted decisions with human review

Lessons Learned
AI-drafted lessons, published by a reviewer

Contradiction detection
Flags documents that disagree with each other

Knowledge-gap detection
Topics people ask about but can't get answered

Topic Evolution
How organizational knowledge shifts over time

Automatic topic discovery
A self-organizing topic structure

Counts toward the EKP score

0

These strengthen the product and our story with NYPA, but none of them closes an EKP requirement. EKP progress is measured only by the 32 requirements.

18 partials: mechanism built, depth missing

Theme	What we have	What's missing
Search & navigation #2–4	Hybrid search, SharePoint file sync, repository filters	Per-file SharePoint permissions, ontology-driven relevance, taxonomy facets
Governance #22–24, #26	Draft → publish → archive → retire; stale-content suggestions	Configurable stages and approvers, version history, owners, audit of changes
Metadata #19, #21	AI summaries; confidence-gated classification	Human review of summaries; required-field rules
Analytics #28–30	Questions, views and knowledge gaps captured	Usage, search-quality and content-health dashboards
Content #11, #13	PDF (with OCR), DOCX, spreadsheets, TXT; in-app viewer	PPTX, images, video; thumbnails
Experience & M365 #1, #5, #10, #32	Dashboard, starred docs, confidence rating, SharePoint read access	Custom landing page, saved searches, explainable scores, SSO

6 requirements have nothing built yet



#12

SharePoint site pages & wikis

Only SharePoint drive files are synced today



#14

Knowledge asset templates

No templates for FAQs, How-To guides or SOPs



#15

Managed metadata

Categories are a fixed list of 7, with no hierarchy or synonyms



#16

Taxonomy-tool integration

No link to NYPA's taxonomy or ontology tool



#25

Configurable review workflows

Review flows are fixed in code



#31

Governance metrics

No reporting on review or validation compliance

Five risks would hurt most in front of NYPA

1	**SharePoint permissions not inherited**	SharePoint is the only mandatory requirement, and our proposal promised permission inheritance. Today access is mapped per site, not per file.
2	**No single sign-on**	Login is email/password or OTP. No Entra ID / M365 sign-in.
3	**Incomplete audit trail**	Document lifecycle changes and deletes aren't logged. Our proposal promised audit history.
4	**No analytics**	We can't yet measure search quality, adoption or time-to-answer.
5	**Fixed taxonomy**	Seven hard-coded categories; no managed metadata or link to NYPA's taxonomy tool.

Six workstreams close all 24 remaining gaps

Workstream	Requirements	Closes
A M365 & connectors (+ single sign-on)	#2, #12, #32	**3**
B Governance & lifecycle	#19, #21–26, #31	**8**
C Taxonomy & semantic layer	#3, #4, #15, #16	**4**
D Analytics & dashboards	#28–30	**3**
E Content coverage	#11, #13, #14	**3**
F Personalization & experience	#1, #5, #10	**3**
Total		**24**

8
met today

+ 24
from the workstreams

= 32
of 32 requirements met

Most of this already sits in PRD Tier 3–4 (approval workflows, audit trail, ownership, personalized workspace). We're pulling it forward, not adding new scope.

Sequencing: biggest risks first, breadth last

1

PHASE 1

Remove the biggest risks

- A · M365 & connectors, single sign-on
- B (part) · Complete the audit trail

SharePoint permissions, SSO, audit trail

2

PHASE 2

Governance & insight

- B · Governance & lifecycle (rest)
- C · Taxonomy & semantic layer
- D · Analytics & dashboards

Configurable workflows, managed metadata, dashboards

3

PHASE 3

Breadth & polish

- E · Content coverage
- F · Personalization & experience

More formats, saved searches, landing page

THANK YOU

A1 · Full requirement status (1 of 4)

#	RFI §3.1 requirement	Status	KMS today / what's missing
1	Customizable landing page	Partial	Org dashboard exists; no admin-configurable landing page
2	Unified search incl. SharePoint (mandatory), Content Server, ServiceNow; respects permissions	Partial	SharePoint file sync; permissions per site, not per file; no Content Server or ServiceNow
3	Hybrid keyword + ontology-driven search	Partial	Vector + BM25 + query and graph expansion; no ontology
4	Faceted navigation on taxonomy	Partial	Flat repository filters; no taxonomy facets
5	Role-based discovery, saved searches, watchlists	Partial	Role-scoped discovery, starred docs; no saved searches or watchlists
6	Recommendations and trending content	Met	Personal recommendations and department trending
7	Natural-language answers alongside results	Met	Org-wide chat; answer box on search page still to add
8	Cited, permission-respecting answers	Met	Typed citations; roles applied before generation

A1 · Full requirement status (2 of 4)

#	RFI §3.1 requirement	Status	KMS today / what's missing
9	Follow-ups and conversational refinement	Met	Conversation and topic context carried forward
10	Confidence scores and explainability	Partial	High/medium/low rating; no numeric score or explanation
11	Formats incl. PPT, video, images	Partial	PDF (OCR), DOCX, spreadsheets, TXT only
12	Site pages, wikis, knowledge articles	Gap	SharePoint drive files only
13	Previews and thumbnails	Partial	In-app viewer; no thumbnails
14	Templates for knowledge asset types	Gap	None
15	Managed metadata: hierarchy, synonyms	Gap	Fixed 7-value category list
16	Integrate with NYPA's taxonomy tool	Gap	None

A1 · Full requirement status (3 of 4)

#	RFI §3.1 requirement	Status	KMS today / what's missing
17	Graph-enriched search, navigation, chat	Met	Knowledge graph used in all three
18	Drag-and-drop upload	Met	Batch upload with category and department
19	Auto description with human review	Partial	AI summaries; no human review step
20	Auto-classification with human review	Met	Confidence thresholds + Needs Review queue
21	Metadata completeness before publishing	Partial	Confidence gate only; no required-field rules
22	Configurable publishing workflow	Partial	Fixed draft → publish state machine
23	Archive, retire, supersede, delete with audit trail	Partial	Archive/retire exist; no supersede, versions or audit log
24	Ownership and stewardship	Partial	Implicit owners; no steward fields

A1 · Full requirement status (4 of 4)

#	RFI §3.1 requirement	Status	KMS today / what's missing
25	Configurable review workflows	Gap	Review flows fixed in code
26	Stale content detection and review triggers	Partial	Stale suggestions via manual script; no schedule or alerts
27	Duplicate and near-duplicate detection	Met	Exact + near-duplicates, plus contradictions
28	Content health dashboard	Partial	Health panel: confidence, conflicts, gaps only
29	Usage and adoption dashboard	Partial	Data captured; no dashboard or search logging
30	Search quality metrics	Partial	Knowledge gaps tracked; no no-results tracking
31	Governance metrics	Gap	None
32	M365 as single source of truth	Partial	Files copied from SharePoint; no SSO or write-back

A2 · Proposal commitments vs. today

Proposal commitment	KMS today	Status
Role-based access and department segmentation	4 roles, enforced in every query	**Met**
Permission inheritance from source systems	Mapped per SharePoint site, not per file (at risk)	**Partial**
Audit of retrieval and AI interactions	Chat audit log, including denied attempts	**Met**
Audit of lifecycle and admin activity	Only org and department events logged (at risk)	**Partial**
Encryption in transit and at rest	Integration tokens encrypted; OpenAI key stored in plaintext	**Partial**
Human approval workflows	Review queues exist but aren't configurable	**Partial**
Secure, source-grounded AI	Permission-filtered retrieval; refuses unsupported answers	**Met**
Pilot measures search quality and time-to-answer	Not instrumented yet (at risk)	**Partial**
Enterprise single sign-on (implied)	Email/password and OTP only	**Gap**

A3 · RFI §4.8 EKP deliverables readiness

§4.8 deliverable	Readiness	Status
Permission-aware unified search integration framework	Connector framework with SharePoint built; others and per-file permissions missing	**Partial**
Semantic/hybrid search pipeline	Hybrid, expansion and graph pipeline built; ontology layer missing	**Partial**
Knowledge graph integration approach	Graph used in search, navigation and chat; ready to document	**Met**
Metadata and taxonomy/ontology mapping specification	Needs managed taxonomy first	**Gap**
Ingestion and enrichment workflows	Pipeline, classification and review queue exist; completeness rules missing	**Partial**
Platform analytics data model	Underlying data captured; no analytics model or dashboards	**Partial**
Retrieval-augmented Q&A architecture	Built; answer-quality evaluation missing	**Partial**